

LlamaIndex with Next.js



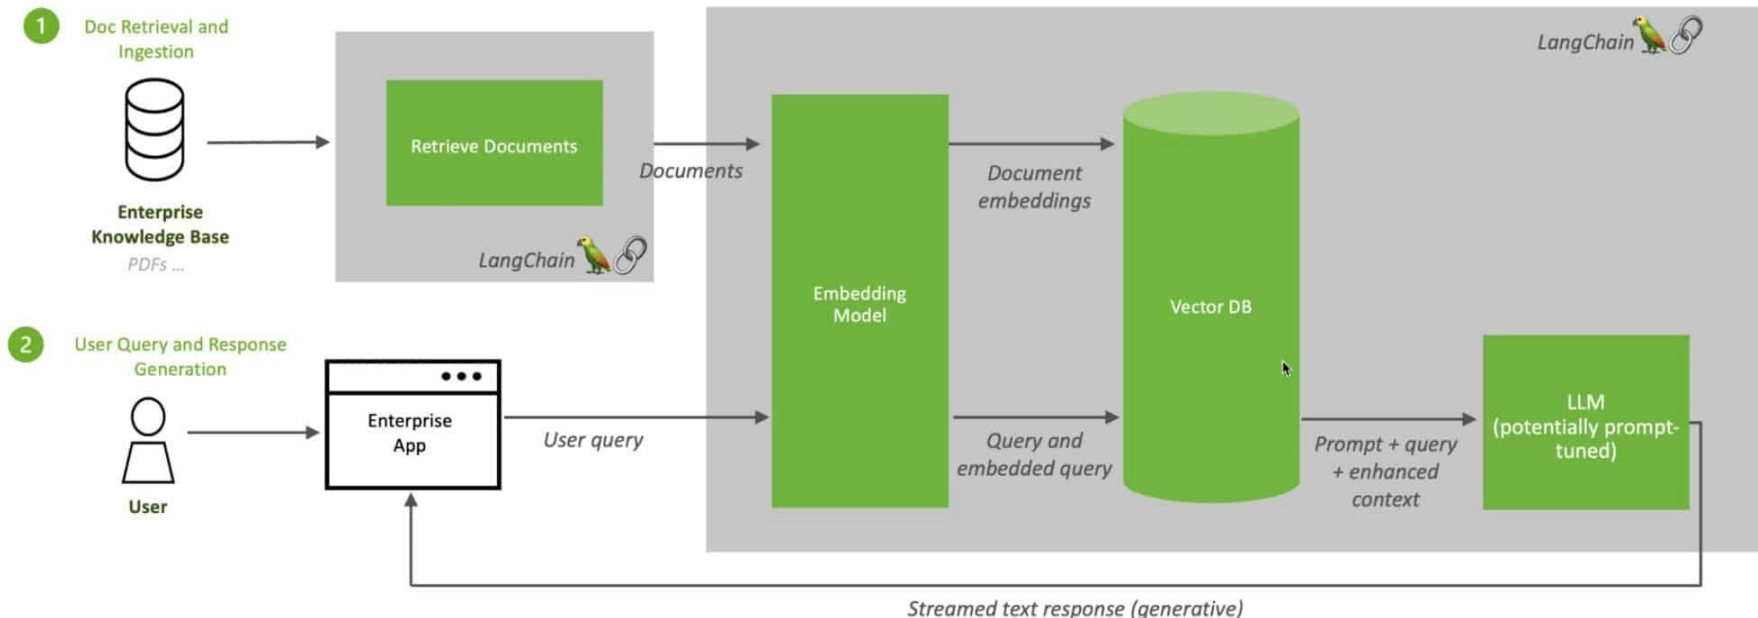
<https://github.com/danielbank/llamaindex-nextjs-demo>

Daniel Bank

A Brief Review of Retrieval Augmented Generation (RAG)


Traditional RAG Pipeline (from [NVIDIA's blog](#))

Retrieval Augmented Generation (RAG) Sequence Diagram




Embedding Models

- An embedding model is a neural network that converts text into dense vector representations (embeddings).
- Different types of Embedding Models based on modality (text, image, audio)
- [Massive Text Embedding Benchmark \(MTEB\)](#)



Vector Similarity (Dot Product)



$$\vec{A} \cdot \hat{B} = |A||B| \cos(\theta)$$

if the magnitude of B is 1, then...

$$C = \vec{A} \cdot \hat{B} = |A| \cos(\theta)$$

Databases (A gentle introduction to Vector DBs)

- Relational Databases
 - When we think of databases, relational databases are usually what we picture
 - Great for Structured Data
 - Data is search by columns

| ISBN | Year | Name | Author |
|------------|------|--------------------------------------|---------------|
| 0767908171 | 2003 | A Short History of Nearly Everything | Bill Bryson |
| 039516611X | 1962 | Silent Spring | Rachel Carson |
| 0374332657 | 1998 | Holes | Louis Sachar |

Vector Databases ([A gentle introduction to Vector DBs](#))

- Vast majority of data on the internet is unstructured (images, text, audio, video)
- Data is searched via content rather than keywords (similarity score, e.g. dot product)
 - Side note: PageRank (1998), the algorithm that powers Google Search, uses eigenvectors to determine page rank (similarity of pages to the query keywords)

| Data UID ¹ | Vector representation |
|-----------------------|---|
| 00000000 | [-0.31, 0.53, -0.18, ..., -0.16, -0.38] |
| 00000001 | [0.58, 0.25, 0.61, ..., -0.03, -0.31] |
| 00000002 | [-0.07, -0.53, -0.02, ..., -0.61, 0.59] |

Tour of the LlamaCloud Products

LlamaParse for Parsing Unstructured Documents

- LlamaParse is a proprietary parsing service, exceptionally adept at transforming PDFs containing complex tables into a neatly organized markdown format:
 - What language to use for OCR?
 - Is there any text at different orientations?
 - Are there special instructions for the Embedding Model?
- Get the data we intend to insert into the Vector DB in a well-defined shape
 - Example: We have text data in different file types (.md, .pdf, .csv, etc.)

LlamaCloud for Managed RAG Pipelines

- LlamaCloud is a managed RAG pipeline running on AWS which also can serve as a managed Vector DB ([Qdrant](#))
- You can organize by:
 - Organizations
 - Projects
 - Indexes (Pipelines)
- Data Sources: Google Drive, SharePoint, S3 Bucket, Box, OneDrive
- Data Sinks: Pinecone, Postgres, Milvus, Azure AI Search
- Embedding Models: Open AI, Cohere, Gemini, HuggingFace, etc.

LlamaReport for Report Generation




Let's build a
RAG-Powered Web App
with LlamaIndex and Next.js



Demo Time

- GitHub Repo: <https://github.com/danielbank/llamaindex-nextjs-demo>



LlamaIndex - <https://github.com/run-llama/LlamaIndexTS>

- `npx create-llama@latest` - Run the CLI tool to bootstrap a web app
- `npm run generate` - Create the Document Embeddings for the knowledge-base data in `./data` (assuming SimpleDirectoryReader example)
- `npm run dev` - Run the web app

Chat Components with [LlamaIndex Chat UI](#)

Batteries-Included <ChatSection />

```
import { ChatSection } from '@llamaindex/chat-ui'

import { useChat } from 'ai/react'

const ChatExample = () => {

  const handler = useChat()

  return <ChatSection handler={handler} />

}
```


Component Composition

```
import { ChatSection, ChatMessages, ChatInput } from '@llamaindex/chat-ui'
import LlamaCloudSelector from '../components/LlamaCloudSelector' // your custom component
import { useChat } from 'ai/react'

const ChatExample = () => {
  const handler = useChat()
  return (
    <ChatSection handler={handler}>
      <ChatMessages />
      <ChatInput>
        <ChatInput.Preview />
        <ChatInput.Form className="bg-lime-500">
          <ChatInput.Field type="textarea" />
          <ChatInput.Upload />
          <LlamaCloudSelector /> { /* custom component */ }
          <ChatInput.Submit />
        </ChatInput.Form>
      </ChatInput>
    </ChatSection>
  )
}
```